

Large-Scale and Multi-Structured Databases

LargeB&B

Grillo Alessio
Sfar Omar
Rotelli Luca

Application Highlights

LargeB&B is a web application that acts as an intermediary platform for booking accommodations. It enables travelers to search, compare, and book hotels, bed & breakfasts, vacation rentals, apartments, and other types of lodging. At the same time, the platform provides property owners and managers with the opportunity to promote their accommodations to a global audience, monitor bookings, and access detailed analytics on the performance of their properties.

The **main features** of LargeB&B include:

- Advanced search and filtering of available rooms based on location and review ratings.
- The ability for users to write and read reviews to help other travelers make informed decisions.
- Real-time chat communication between customers and property managers, both before and after booking.
- Access to analytics and performance insights for property owners regarding the accommodations they have published.

Actors (i)

Customer and unregistered user

The screenshot shows a web interface for a property search platform. At the top, there is a header with the logo "LargeB&B" and a search bar containing the text "From Jan 25, 2026 to Jan 28, 2026...". Below the header is a "Filters" sidebar on the left. The sidebar includes dropdown menus for "Location" (Country and Region) and a section for "Services" with checkboxes for WiFi, Pool, Parking, Gym, and Air Conditioning. It also indicates "From \$150 reviews". To the right of the sidebar, the main content area is titled "Properties". It features a heading "See all" and three property cards: "Cozy Cottage" (a wooden cabin), "Modern Apartment" (a modern building), and "Historic Villa" (a large, classical villa with a pool). Each card includes a thumbnail image, the property name, its rating (4.7/★★★), and the number of reviews (150, 80, or price per night).

Page where the user search for properties and filter them.

Actors (ii)

Customer and unregistered user

LargeB&B Jam 26, 26, 2026 to Jan 28, 2026...  **Search**

Filters

Location

Country: Tuscania

Region: Tuscania

From \$150 reviews

Services

Breakfast Included

Pet-Friendly

Pool Access

Price Per Night

Adults: \$19000

Children: \$500

Rooms in Florence Sort by: Recommended

 **Deluxe King Room**
4.7/★★★ (150 reviews)

 **Standard Room**
4.7/★★★ (150 reviews)

 **Standard Double**
4.5/★★★ (Private Balcony)
Ensuite Bathroom
Book Now

 **Standard Double**
4.5/★★★ (80 reviews)
Ensuite Bathroom
Book Now

Page where the user search for rooms and filter them.

Actors (iii)

Customer and unregistered user

The screenshot shows a web page for a travel booking platform. At the top, there is a search bar with the placeholder "Cerca una località..." and a magnifying glass icon. To the right of the search bar are buttons for "Becomes" and "Host" with a user profile icon.

The main content area features a listing for a room named "Astro" with a rating of 4.98. The listing includes a title "Sunny Loft with Private Terrace near Colosseum", a large image of the terrace with outdoor furniture and potted plants, and two smaller images: one of the terrace and another of a bathroom with a bathtub.

Below the image, there is a descriptive text block: "A bright refuge è the heart in the Eternal City". It continues: "Immerse as prógue in the Dolce Vita' in this spéndid loft this recently desally clmt reçoett Roman charia aue true xuâte rerorate panoramie. Punshed deyerrace with hzinished terrace withn the cupriqianc plantss. High-speed-WFi and and smart working atta station."

On the left side of the listing, there is a "Reviews" section with an average rating of ★ 4.98 (124 reviews). It includes a review from Giulia M. (October 2025) and another from Mark & Sarh (September 2025).

On the right side, there is a "Pricing per night:" section with rates for adults and children, an estimated total of €290.00 + cleaning fees, and an estimated total of €145.00 for two nights. A "BOOK NOW" button is available, along with a note that you won't be charged at stage.

Page where the user can see the room details.

Actors (iv)

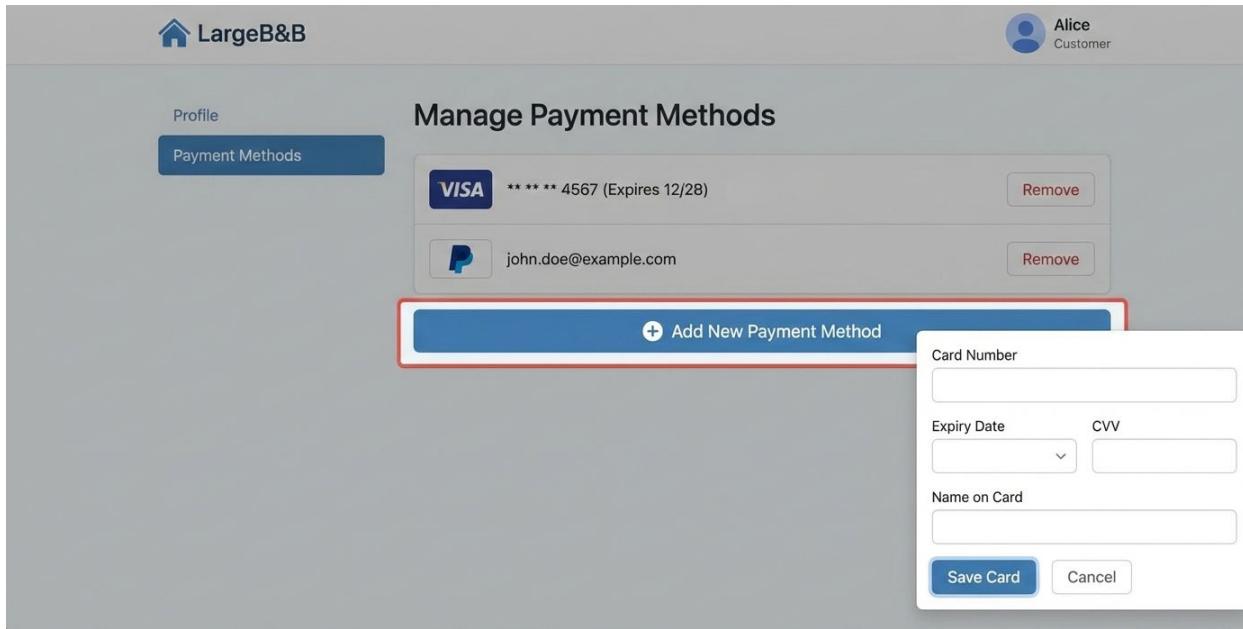
Customer

The screenshot shows a web application interface for managing bookings. At the top, there is a dark header bar with the logo 'LargeB&B' and navigation links: Dashboard, Discover, Account, Astro, and a user icon. Below the header, the main title 'Your Bookings' is displayed. Underneath, a section titled 'Upcoming' lists a booking for 'Sunny Loft Colosseum' in Rome, Italy, from March 15 to March 19, 2026. It includes a thumbnail image of a terrace, the booking ID '2026-3-15-A', and buttons for '[Manage Booking]', 'Add to Calendar', and a gear icon. To the right of this listing, the total cost is shown as 'Booked total: 115.65'. Below this, another section titled 'Past Bookings' lists two entries: one for 'Historic Villa Amaflili Coast' in August 2025 and another for May 2023. Each past booking entry includes a thumbnail image, the booking date, a 'View Details' button, a 'Leave a Review' button, and a red 'X' button.

Page where the customer can see his reservations.

Actors (v)

Customer



Page where the customer can manage his payment method.

Actors (vi)

Customer

The screenshot shows a web application interface for 'LargeB&B'. At the top, there is a navigation bar with tabs: Dashboard, Bookings, Properties, Properties (selected), and Messages. On the right side of the header is a user profile for 'Alice Customer' with a small profile icon.

The main content area is titled 'My Favorite Properties'. On the left, there is a sidebar with 'Profile' and 'Favorites' sections. The 'Favorites' section is highlighted with a blue background and contains a button labeled 'Favorites'.

The main content area displays a grid of six property cards:

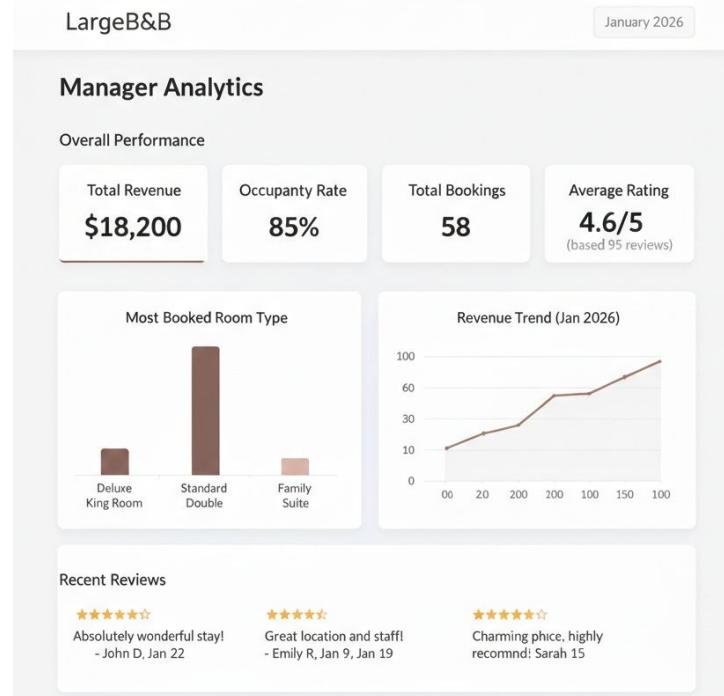
- Seaside Villa**
Malibu, California
4 beds, 3 baths
- Mountain Retreat**
Aspen, Colorado
3 beds, 2 baths
- Mountain Retreat**
Aspen, 2 baths
2 beds, 1 bath
- Urban Loft**
- Urban Loft**
- Urban Loft**

Each property card features a small heart icon in the top right corner, indicating it is a favorite. A 'Remove from Favorites' button is located in the top right corner of the main content area.

Page where the customer can see his favorite properties.

Actors (vii)

Manager



Page where the manager can see his analytics.

Actors (viii)

Manager

The screenshot shows the LargeB&B Manager dashboard titled "Manage Your Properties". On the left, there's a list of properties: "Sumy Loft Colossem", "Historic Villa Amali Coast", and "Urban Studio Paris". Each property card includes a thumbnail image, the property name, a "View Property" button, and an "Edit Property" button. To the right of the properties is a section titled "Actions" containing buttons for "Add/Remove a Room" and "View All Bookings". Below these actions is a modal window titled "Edit Room Details" with fields for "Select Property" (set to Sumy Loft Colossum), "Select Room" (set to Master Bedroom Terrace), "Capacity Nite" (set to Five pes Night), and "Price per Night" (set to Upload Photos). There are also "Upload Photos", "Save", and "Cancel" buttons at the bottom of the modal.

Page where the manager can manage his properties.

Actors (ix)

Manager and Customer

The screenshot shows a web-based messaging interface for a vacation rental platform. At the top, there's a navigation bar with tabs: Dashboard, Bookings, Properties, and Messages. The 'Messages' tab is currently selected. In the top right corner, there's a user profile for 'Alice Customer'. On the left, there's a sidebar titled 'Conversations' with three entries: 'Manager: John Doe (Seaside Villa)' (with a message 'Looking forward to your stay!'), 'Manager: Sarah Smith (Mountain Retreat)' (with a message 'Looking forward to... Mountain retreat!'), and 'Support Team' (with a message 'We've sent to you for a property mana...'). The main area is a chat window titled 'Chat with Manager: John Doe - Booking #12345 (Seaside Villa, Aug 12-15)'. It shows a back-and-forth conversation between Alice Customer and John Doe. Alice asks about arrival time and parking, and John confirms both. The interface includes a text input field at the bottom with placeholder 'Type your message here...', a 'Send' button, and some small icons.

LargeB&B

Dashboard Bookings Properties Messages

Alice Customer

Conversations

Manager: John Doe (Seaside Villa)
Looking forward to your stay!

Manager: Sarah Smith (Mountain Retreat)
Looking forward to... Mountain retreat!

Support Team
We've sent to you for a property mana...

Chat with Manager: John Doe - Booking #12345 (Seaside Villa, Aug 12-15)

[Today, 10:05 AM]
John Doe: Hi Alice, thank you for your booking! We're excited to host you at Seaside Villa.

[Today, 10:06 AM]
John Doe: Just confirming, your estimated arrival time is around 3 PM, correct?

[Today, 10:15 AM]
Alice Customer: Hello John! Yes, that's correct. We should arrive by 3 PM.

[Today, 10:16 AM]
Alice Customer: Also, could you confirm if there's parking available on-site? Thanks!

[Today, 10:20 AM]
John Doe: Yes, we have a dedicated parking space right in front of the villa for you. See you soon!

Type your message here... Send 📎 😊

Page where the manager or customer can send/receive messages.

Actors (x)

Manager and Customer

The screenshot shows a web application interface for 'LargeB&B'. At the top, there is a navigation bar with links: Dashboard, Bookings, Properties, Properties (selected), and Messages. On the right side of the header, there is a user profile icon labeled 'Alice Customer'.

The main content area is titled 'Notifications'. It displays four notification cards:

- New Message:** Mario has written to you
you
Today, 9:55 PM
- New Booking Made:** Customer Luigi has new reservation for Seaside Villa
Today, 9:30 PM
- Booking Cancelled:** Customer Fabio has cancelled their reservation for Mountain Retreat
Yesterday, 4:15 PM
- Booking Modified:** Customer Gianni updated the reservation details for Urban Loft
Yesterday, 10:00 AM

Page where the manager or customer can see his notifications (booking notification are only for Manager).

Dataset Description

Sources: Inside AirBnB, Faker

Description: Dataset concerning properties and rooms information, reviews, notifications, messages between customer and manager, and reservations information.

Volume: 170MB

Variety: Frequently changing data: real-time room availability, new bookings, and daily reviews.

Dataset Description

Velocity/Variability:

The system handles detailed **room and property descriptions**, along with essential booking data such as **guest counts** (adults and children) and **check-in/check-out dates**. It features a comprehensive review system that includes **multi-criteria ratings** for stays and allows for **manager responses**. Additionally, a real-time notification system alerts users to **incoming messages** and provides managers with exclusive updates regarding the **reservation status** of their properties.

Pre-processing:

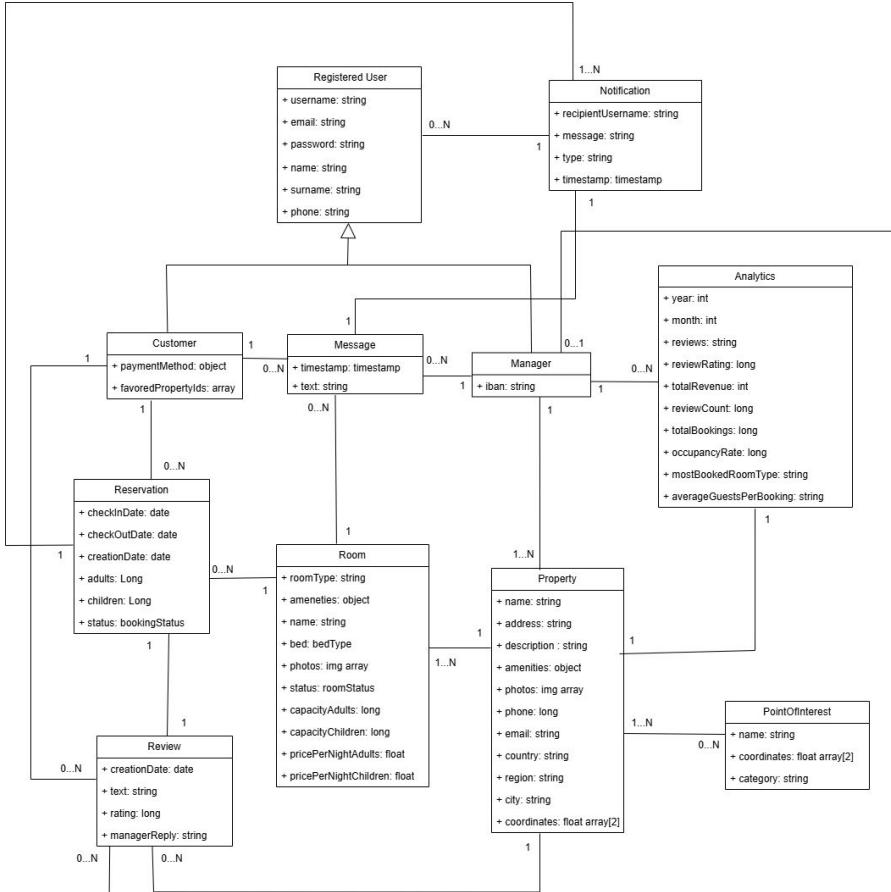
Data Cleaning: Used Python (Pandas) to clean InsideAirbnb CSVs (handling missing values, removing inactive listings).

Data Fusion: Merging real property data with synthetic Manager profiles.

Transformation: Converting flat CSV records into hierarchical MongoDB Documents (Embedding Rooms/Reviews).

Geo-Coding: Transforming lat/con columns into GeoJSON points objects for spatial indexing.

UML Design Class Diagram



Application non-functional requirements

- *Passwords of registered users must be securely encrypted.*
- *The system must provide notifications via application.*
- *The results produced by the system must be consistent.*
- *The system must ensure high data availability.*
- *The system must manage data storage using Document DB, Graph DB, and Key-Value DB.*
- *The system should be able to scale to accommodate an increase in user traffic.*

Application non-functional requirements

- *The DBMS technologies used must be Neo4j, MongoDB, and Redis.*
- *The business logic and data access layers must be implemented using the Spring framework in Java.*
- *The codebase must be easily maintainable, readable, and updatable.*
- *The system must use a document database to store information permanently.*
- *The system must use a key-value database to temporarily store information, manage cache, locks, and user sessions.*

Document DB Design

Properties

```
_id: "bdd640fb-0667-4ad1-9c80-317fa3b1799d"
name : "Elif's room in cozy, clean flat."
address : "VIII Appia Antica"
description : "10 min by bus you can get to Piazza Venezia or Colosseum. All shops, g..."
▶ amenities : Array (42)
▶ photos : Array (5)
email : "williamdavis@callahan.com"
country : "Italy"
region : "Lazio"
city : "Rome"
▶ ratingStats : Object
▶ latestReviews : Array (10)
managerId : "23b8c1e9-3924-46de-beb1-3b9046685257"
▶ location : Object
▶ rooms : Array (3)
▶ pois : Array (6)
```

Reservations

```
_id: "1825bc54-30be-445f-a835-14f2ceb81f9d"
adults : 2
children : 1
status : "confirmed"
userId : "14296c07-f26b-4776-913e-4de2e0c53cb8"
roomId : "ce88cb2d-d4e8-4839-bc3e-058be0f3eab0"
▶ dates : Object
createdAt : 2025-11-19T00:00:00.000+00:00
```

Rooms (embedded)

```
roomType : "single"
▶ amenities : Array (42)
name : "Room 1"
beds : 2
▶ photos : Array (3)
status : "available"
capacityAdults : 3
capacityChildren : 2
pricePerNightAdults : 92.25
pricePerNightChildren : 19.42
roomId : "e2817efd-ae84-4217-9d53-434bb88139b9"
```

Pois (embedded)

```
id : "113634a5-2050-4c62-88e8-95d7702c1821"
name : "Colosseum"
coordinates : Array (2)
type : "historical"
```

Messages

```
_id: "ee87905e-4ca4-45ea-8dfa-6a56d12dbc9a"
timestamp : 2025-11-24T00:56:00.000+00:00
content : "Voice care break blood network evening painting."
senderId : "da4bd9ca-eb5c-4467-80ba-cd647a0ecfea"
recipientId : "23b8c1e9-3924-46de-beb1-3b9046685257"
isRead : true
```



Document DB Design

Reviews

```
_id: "0b19f88e-9d77-445e-b206-c26938b77c07"  
propertyId : "bdd640fb-0667-4ad1-9c80-317fa3b1799d"  
userId : "6712303a-0f84-4fef-9931-e9eea56c0941"  
reservationId : "0ef8c2d6-f7fd-4646-b7bb-3eec4bf50b52"  
creationDate : 2014-12-26T00:00:00.000+00:00  
text : "This spot was in a great and nice area, walking a distance about 700 m:..."  
rating : 5  
cleanliness : 3  
communication : 5  
location : 4  
value : 3  
managerReply : "Feeling poor all your suggest international blue."
```

Notifications

```
_id: "6c6fa611-5ab3-4edf-ae59-5ed3a8b317fa"  
recipientId : "23b8c1e9-3924-46de-beb1-3b9046685257"  
title : "New Booking"  
body : "Customer Kyle has made a new reservation."  
type : "RESERVATION_CREATED"  
referenceId : "1825bc54-30be-445f-a835-14f2ceb81f9d"  
read : false  
timestamp : 2025-11-19T00:00:00.000+00:00
```

Users (Customer)

```
_id: "14296c07-f26b-4776-913e-4de2e0c53cb8"  
username : "dennislisa"  
email : "davidalvarez@example.net"  
password : "#4RQ2tzaKQ"  
name : "Kyle"  
surname : "Brown"  
birthdate : 1966-05-05T00:00:00.000+00:00  
phoneNumber : "001-330-510-3105"  
► paymentMethod : Object  
► favoredPropertyIds : Array (empty)  
role : "CUSTOMER"
```

Users (Manager)

```
_id: "23b8c1e9-3924-46de-beb1-3b9046685257"  
username : "rhodespatricia"  
email : "garzaanthony@example.org"  
password : "@2$3Qt!ld"  
name : "Susan"  
surname : "Rogers"  
phoneNumber : "361-855-9407"  
iban : "GB38CRJU95931034131647"  
vatNumber : "KD52553419283"  
► billingAddress : Object  
role : "MANAGER"
```

Document DB Design

Aggregation Queries Benefiting from Design

```
// Advanced Search - amenities embedded enables $all operator
db.properties.aggregate([
  { $match: { amenities: { $all: ["Wifi", "Pool"] }, city: "Rome" } },
  { $sort: { "ratingStats.value": -1 } }
])

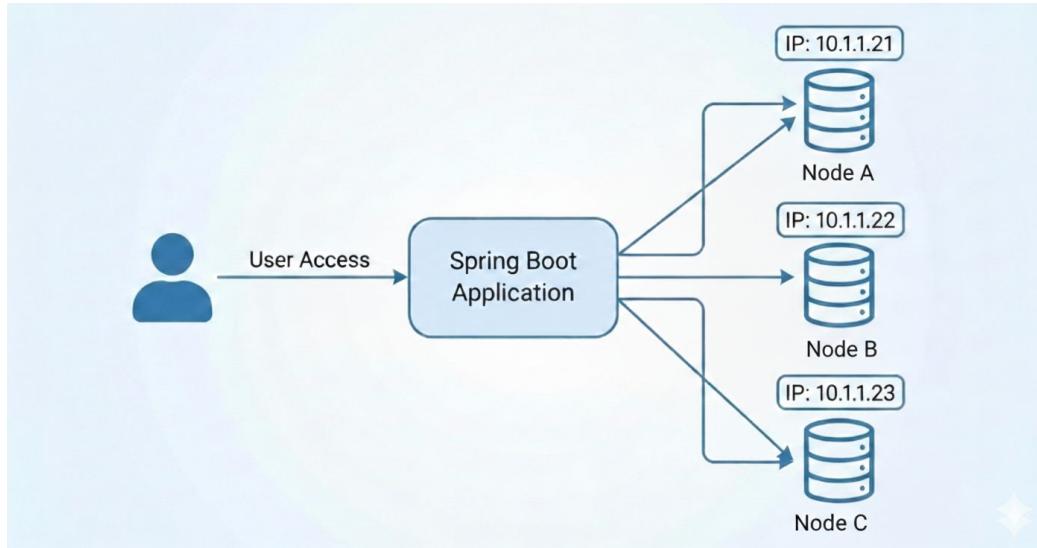
// Geospatial Query - coordinates embedded in Property
db.properties.aggregate([
  { $geoNear: { near: { type: "Point", coordinates: [12.4964, 41.9028] },
                distanceField: "distance", maxDistance: 5000, spherical: true } }
])

// Top Rated Properties - ratingStats pre-computed
db.properties.aggregate([
  { $match: { "ratingStats.count": { $gte: 10 } } },
  { $addFields: { avgRating: { $divide: ["$ratingStats.sum", "$ratingStats.count"] } } },
  { $sort: { avgRating: -1 } }, { $limit: 10 }
])
```

We adopted an embedded model for data always queried together (Property+Rooms), and a referenced model for independent but correlated entities (Reservation ← Room). This enables fast queries such as amenity-based search, geolocation, and rating-based sorting without expensive JOINs. Each application view queries data exactly as it is structured, maximizing read performance.

Cluster Representation

The system deployment involves a 3-node cluster configured to maximize resource efficiency and availability.



MongoDB Replica Set Configuration

We deployed a **MongoDB Replica Set** across all three nodes with a priority-based election strategy: Node A acts as the **Primary** (Priority 5), while Nodes B (Priority 2) and C (Priority 1) serve as Secondaries.

MongoDB Replica Set Configuration

Read Operations (Read Preference: nearest)

Given that the read-to-write ratio in LargeB&B is estimated to be approximately 10:1 (browsing vs. booking), minimizing read latency is crucial. We utilize the nearest read preference. The driver routes the read request to the replica set member with the lowest network latency, regardless of whether it is Primary or Secondary.

This choice provides two main benefits:

- Load Balancing: It distributes the heavy read traffic across all three nodes, preventing the Primary from being overwhelmed.
- Geo-latency Reduction: If the infrastructure is geo-distributed in the future, users will automatically read from the data center closest to them.

MongoDB Replica Set Configuration

Write Operations (Write Concern: majority):

Although we prioritize availability for reads, write operations—specifically User Registration, Property Creation, and critical Reservation flows—require strict durability. We have configured the Write Concern to majority.

- The system acknowledges a write only after it has been committed to the Primary and propagated to at least one Secondary (2 out of 3 nodes).
- This ensures that even if the Primary node crashes immediately after a write (e.g., after a confirmed payment), the data is safe on a Secondary node that will become the new Primary, preventing data rollback and financial inconsistencies.

MongoDB Indexes

Left image query example with no index & right image query example with index

Properties Collection:

`{"rooms.capacityAdults": 1, "rooms.capacityChildren": 1}`: It allows the database to efficiently filter properties that contain at least one room capable of accommodating the requested number of adults and children.

```
executionStats: {  
    executionSuccess: true,  
    nReturned: 1920,  
    executionTimeMillis: 119,  
    totalKeysExamined: 0,  
    totalDocsExamined: 1950,
```

```
executionStats: {  
    executionSuccess: true,  
    nReturned: 1920,  
    executionTimeMillis: 85,  
    totalKeysExamined: 5337,  
    totalDocsExamined: 1933,
```

MongoDB Indexes

"ratingStats.value": Optimizes the "Sort by Highest Rated" feature, ensuring that high-rated properties are retrieved efficiently without performing expensive in-memory sorting

```
executionStats: {  
    executionSuccess: true,  
    nReturned: 1950,  
    executionTimeMillis: 178,  
    totalKeysExamined: 0,  
    totalDocsExamined: 1950,
```

```
executionStats: {  
    executionSuccess: true,  
    nReturned: 1950,  
    executionTimeMillis: 45,  
    totalKeysExamined: 1950,  
    totalDocsExamined: 1950,
```

MongoDB Indexes

"managerId": Ensures fast retrieval of the specific subset of properties owned by a logged-in manager for the dashboard view

```
executionStats: {  
    executionSuccess: true,  
    nReturned: 1,  
    executionTimeMillis: 5,  
    totalKeysExamined: 0,  
    totalDocsExamined: 1950,
```

```
executionStats: {  
    executionSuccess: true,  
    nReturned: 1,  
    executionTimeMillis: 0,  
    totalKeysExamined: 1,  
    totalDocsExamined: 1,
```

MongoDB Indexes

Text-Index:

*"name (property) (weight: 3), name (room) (weight: 2), description (weight: 1), city (weight: 1)":
Enables full-text search capabilities with relevance scoring. Matches in the property name are prioritized over matches in the description, improving search relevance for users*

```
executionStats: {  
    executionSuccess: true,  
    nReturned: 354,  
    executionTimeMillis: 148,  
    totalKeysExamined: 406,  
    totalDocsExamined: 354,  
}
```

MongoDB Indexes

2dsphere Index:

"coordinates": Required to support geospatial queries such as \$near and \$geoWithin, satisfying the requirement to view properties on a map or find POIs nearby.

```
executionStats: {  
    executionSuccess: true,  
    nReturned: 135,  
    executionTimeMillis: 36,  
    totalKeysExamined: 165,  
    totalDocsExamined: 139,
```

MongoDB Indexes

Rooms (Embedded in Property collection):

“pricePerNightAdults”: Optimizes complex budget queries where users filter by adult price

```
executionStats: {  
    executionSuccess: true,  
    nReturned: 1,  
    executionTimeMillis: 5,  
    totalKeysExamined: 0,  
    totalDocsExamined: 1950,
```

```
executionStats: {  
    executionSuccess: true,  
    nReturned: 1,  
    executionTimeMillis: 0,  
    totalKeysExamined: 1,  
    totalDocsExamined: 1,
```

MongoDB Indexes

Reservations Collection:

`{ "roomId": 1, "dates.checkIn": 1, "dates.checkOut": 1 }` : It allows the system to efficiently check if a specific room is occupied during a given date range before confirming a new booking

```
executionStats: {  
    executionSuccess: true,  
    nReturned: 1,  
    executionTimeMillis: 34,  
    totalKeysExamined: 0,  
    totalDocsExamined: 29644,
```

```
executionStats: {  
    executionSuccess: true,  
    nReturned: 1,  
    executionTimeMillis: 0,  
    totalKeysExamined: 1,  
    totalDocsExamined: 1,
```

MongoDB Indexes

"userId": Optimizes the retrieval of booking history for the "My Bookings" page

```
executionStats: {  
    executionSuccess: true,  
    nReturned: 1,  
    executionTimeMillis: 859,  
    totalKeysExamined: 0,  
    totalDocsExamined: 29644,
```

```
executionStats: {  
    executionSuccess: true,  
    nReturned: 1,  
    executionTimeMillis: 77,  
    totalKeysExamined: 1,  
    totalDocsExamined: 1,
```

MongoDB Indexes

"userId": Allows the system to quickly retrieve all reservations associated with a specific room

```
executionStats: {  
    executionSuccess: true,  
    nReturned: 1,  
    executionTimeMillis: 859,  
    totalKeysExamined: 0,  
    totalDocsExamined: 29644,
```

```
executionStats: {  
    executionSuccess: true,  
    nReturned: 1,  
    executionTimeMillis: 77,  
    totalKeysExamined: 1,  
    totalDocsExamined: 1,
```

Discussion on MongoDB Data Sharding

*Since the **id** field is mandatory in every document and is the primary filter for the majority of queries, it was selected as the shard key. This choice enables direct query routing, avoiding the overhead of scatter-gather operations. Consequently, we adopted Hashed Sharding as the partitioning algorithm to ensure uniform distribution.*

Key-value DB Design

Redis Data Model & Key Naming

Key Naming Strategy:

1. JWT Blacklist: "blacklist:{token}" → Value: "true" (String, TTL auto-expire)
2. Pending Reservations: "reservation:temp:{uuid}" → Value: Reservation JSON (15min TTL)
3. Trending Properties: "trending_properties" → Sorted Set (propertyId, score)
4. Recently Viewed: "recent:{userId}" → List of propertyIds (max 10)

Design Justification:

- **Namespace prefixes** (blacklist:, reservation:temp:, recent:) prevent key collisions
- **TTL expiration** for temporary data (JWT tokens, pending reservations) ensures automatic cleanup
- **Sorted Set** for trending enables O(log N) scoring and range queries
- **List** for recently viewed maintains insertion order with O(1) head insertion

Key-value DB Design

Implemented Operations

```
// 1. JWT Token Blacklisting (Secure Logout)
redisTemplate.opsForValue().set(
    "blacklist:" + token, "true", ttl, TimeUnit.MILLISECONDS
);
Boolean isBlacklisted = redisTemplate.hasKey("blacklist:" + token);

// 2. Pending Reservation (Temporary Cart)
redisTemplate.opsForValue().set(
    "reservation:temp:" + uuid, reservation, 15, TimeUnit.MINUTES
);

// 3. Trending Properties (Real-time Analytics)
redisTemplate.opsForZSet().incrementScore("trending_properties", propertyId, 1);
Set<Object> topIds = redisTemplate.opsForZSet().reverseRange("trending_properties", 0, 9);

// 4. Recently Viewed History
redisTemplate.opsForList().leftPush("history:" + userId, propertyId);
redisTemplate.opsForList().trim("history:" + userId, 0, 9); // Keep only last 10
```

Key-value DB Design

System Benefits:

- **Security:** Stateless JWT invalidation without database queries
- **Performance:** Sub-millisecond access for trending properties (vs MongoDB aggregation)
- **UX:** Pending reservations survive page refresh during payment flow
- **Personalization:** Recently viewed history without polluting main DB

Views Supported:

- **ShowProperties** - Displays trending properties banner (reads `trending_properties` ZSet)
- **Checkout Flow** - Retrieves pending reservation during payment (`reservation:temp:{id}`)
- **User Dashboard** - Shows recently viewed properties (`recent:{userId}` List)

Redis Replica Set Configuration

*For Redis, Node A acts as the **Master** handling both write and read operations, supported by two Read-only Replicas on Nodes B and C. High Availability is guaranteed by **3 Redis Sentinels**, which continuously monitor the cluster's health and automatically promote a Replica to Master in the event of a failure. Data durability and disaster recovery are secured through periodic **RDB snapshots** persisted to disk.*

*To ensure system stability under heavy load, we implemented a strict memory management strategy. We set a maxmemory **limit of 512MB** combined with a **Least Frequently Used (LFU)** eviction policy. When this limit is reached, Redis evicts the least accessed keys first, protecting the system from memory saturation and crashes during traffic spikes.*

Discussion on Redis Data Sharding

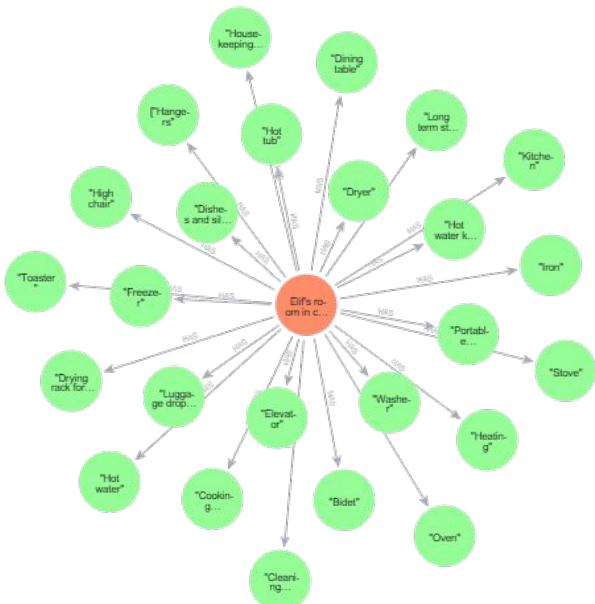
We implemented Sharding in Redis for:

- **Horizontal Scalability (RAM):** It allows storing datasets larger than the memory limit of a single server by distributing data across multiple nodes.
- **Increased Write Throughput:** Since Redis is single-threaded, sharding splits write operations across multiple CPUs, significantly increasing total cluster performance.

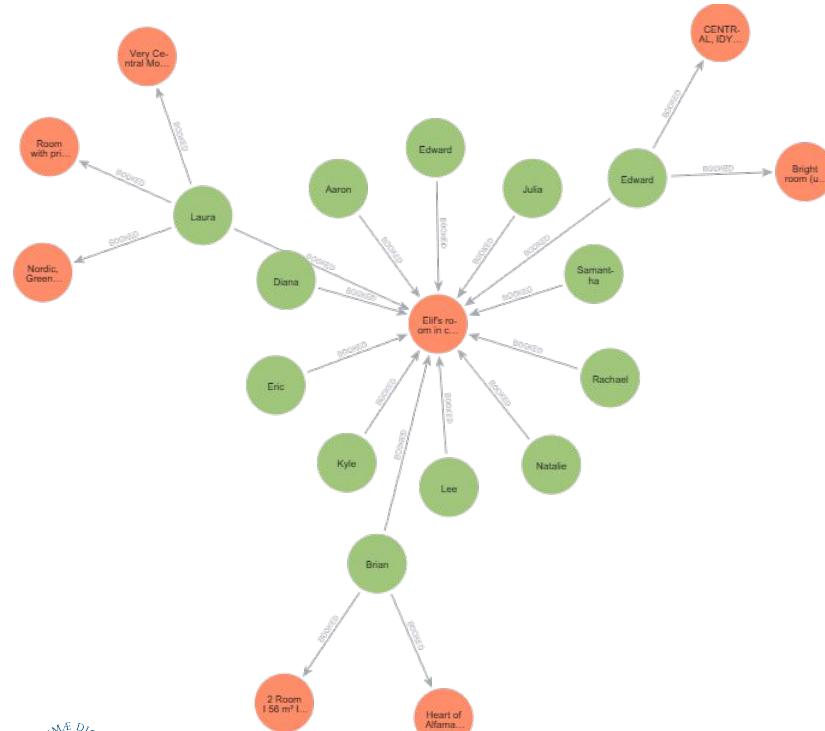
*The database key itself serves as the **sharding key**, and the partitioning strategy relies on a **hashing algorithm**.*

Graph DB Design

(:Property)-[:HAS]->(:Amenity)



(:User)-[:BOOKED]->(:Property)



Graph DB Design

Collaborative Filtering (User-Based Recommendations)

```
// Find properties booked by users with similar booking patterns
MATCH (p:Property {propertyId: $propId})-[:BOOKED]-(u:User)
      -[:BOOKED]->(other:Property)
RETURN other.propertyId AS recommendedId,
       count(*) AS strength
ORDER BY strength DESC
LIMIT 5
```

Content-Based Filtering (Feature Similarity)

```
// Find properties sharing common amenities
MATCH (p:Property {propertyId: $propId})-[:HAS]->(a:Amenity)
      <-[ :HAS ]-(other:Property)
WHERE p.propertyId <> other.propertyId
RETURN other.propertyId AS recommendedId,
       count(a) AS matchStrength
ORDER BY matchStrength DESC
LIMIT 10
```

Graph DB Design

Performance:

- Graph traversal in $O(\text{depth} \times \text{branching_factor})$, independent of total dataset size
- 2-hop collaborative filtering: ~10ms vs ~500ms MongoDB aggregation

Query Expressiveness:

- "Users who booked X also booked Y" translates naturally to graph pattern matching
- Amenity-based similarity requires no complex array operations

Decoupled Architecture:

- Neo4j stores projection → can be rebuilt from MongoDB if corrupted
- Read-only for recommendations → no write conflicts with transactional data

Views Supported:

- **Property Details Page** → Shows "Similar Properties" (content-based query)
- **Recommendations Section** → Shows "Users Also Booked" (collaborative filtering)

Neo4j Indexes

Left image query example with no index & right image query example with index

Property Node:

"PropertyId": Both recommendation algorithms begin with a specific property currently being viewed (the "Anchor Node"). This constraint ensures O(1) lookup time to locate this starting point in the graph immediately

Cypher version: 5

Planner: COST

Runtime: PIPELINED

Total database accesses: 3.907

Total allocated memory: 312 bytes

Total time: 467 ms

Cypher version: 5

Planner: COST

Runtime: PIPELINED

Total database accesses: 4

Total allocated memory: 312 bytes

Total time: 363 ms

Neo4j Indexes

User Node:

“UserId”: Essential for the Collaborative Filtering algorithm

Cypher version: 5

Planner: COST

Runtime: PIPELINED

Total database accesses: 63.195

Total allocated memory: 312 bytes

Total time: 283 ms

Cypher version: 5

Planner: COST

Runtime: PIPELINED

Total database accesses: 4

Total allocated memory: 312 bytes

Total time: 276 ms

Neo4j Indexes

Amenity Node:

"AmenyName": Ensures that shared features (e.g., "WiFi", "Pool") are represented as single, unique nodes in the graph. This allows the traversal algorithm to "hop" from the current property to related properties via these shared amenity nodes efficiently

Cypher version: 5

Planner: COST

Runtime: PIPELINED

Total database accesses: 6

Total allocated memory: 312 bytes

Total time: 286 ms

Cypher version: 5

Planner: COST

Runtime: PIPELINED

Total database accesses: 3

Total allocated memory: 312 bytes

Total time: 126 ms

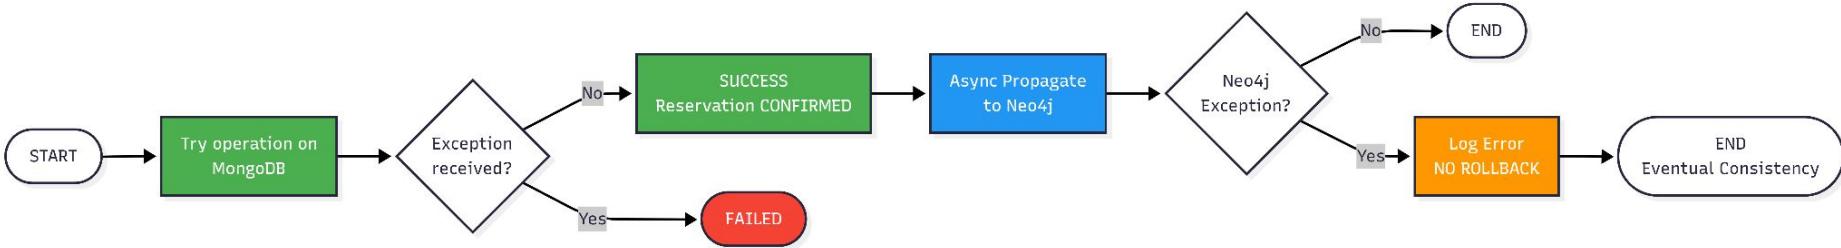
Handling Intra-DB Consistency

Hypothesis: Since MongoDB has 3 replicas with Write Concern "majority", we can assume that a fail on connection or data persistence is almost impossible.

This consistency is managed using **asynchronous propagation** with fire-and-forget pattern:

- **MongoDB is the Source of Truth:** All critical write operations (reservations, payments) are persisted to MongoDB FIRST with strong consistency guarantees.
- **Neo4j is a Projection:** After MongoDB confirms the transaction, the relationship is propagated to Neo4j asynchronously using CompletableFuture.
- **Eventual Consistency Model:** If Neo4j propagation fails, the system logs the error but does NOT rollback the MongoDB transaction. The reservation remains valid.
- **Acceptable Inconsistency Window:** There may be a brief delay (milliseconds) before the booking appears in the recommendation graph, but this doesn't impact transactional integrity or user experience.

Handling Intra-DB Consistency



```
CompletableFuture.runAsync(() -> {
    try {
        reservationGraphRepository.createReservation(
            finalReservationId,
            finalUserId,
            finalPropertyId,
            finalReservation.getDates().getCheckIn(),
            finalReservation.getDates().getCheckOut(),
            finalAmount
        );
        logger.info("Graph: Reservation {} synced to Neo4j successfully.", finalReservationId);
    } catch (Exception e) {
        logger.error("Graph Error: Failed to sync reservation {} to Neo4j: {}", finalReservationId, e.getMessage());
    }
});
```

Swagger UI REST APIs documentation

Swagger doc url: <http://localhost:8080/swagger-ui/index.html>

property-controller

GET /api/properties/{propertyId}

GET /api/properties/{propertyId}/recommendations/similar

GET /api/properties/{propertyId}/recommendations/collaborative

GET /api/properties/trending

GET /api/properties/top-rated

GET /api/properties/search

GET /api/properties/map

reservation-controller

PUT /api/reservations/{reservationId}

DELETE /api/reservations/{reservationId}

POST /api/reservations/initiate

POST /api/reservations/confirm-payment

GET /api/reservations/my-reservations

DELETE /api/reservations/temp/{tempReservationId}

Live Demo with Postman

The screenshot shows the Postman interface with the 'Large B&B' collection selected. The left sidebar displays a tree view of API endpoints under 'api': reviews, {reviewId}, reply, property, reservations, manager, payment-methods, messages, favored, auth, notifications, properties, and users. The 'reviews' endpoint has three sub-endpoints: PATCH reply To Review, PUT edit Review, and DEL delete Review. The 'reply' endpoint has one sub-endpoint: POST create Review. The 'property' endpoint has one sub-endpoint: POST create Review. The 'reservations' endpoint has one sub-endpoint: POST create Reservation. The 'manager' endpoint has one sub-endpoint: POST create Manager. The 'payment-methods' endpoint has one sub-endpoint: POST create Payment Method. The 'messages' endpoint has one sub-endpoint: POST create Message. The 'favored' endpoint has one sub-endpoint: POST create Favored. The 'auth' endpoint has one sub-endpoint: POST create Auth. The 'notifications' endpoint has one sub-endpoint: POST create Notification. The 'properties' endpoint has one sub-endpoint: POST create Property. The 'users' endpoint has one sub-endpoint: POST create User. The main panel shows the title 'Large B&B' and a descriptive text: 'This application is a comprehensive platform for managing room bookings and property listings. It leverages a polyglot persistence architecture, integrating a MongoDB cluster for high-availability document storage, Neo4j for relationship-based recommendations, and Redis for distributed locking and concurrency management. Have a great stay!'. Below the text is a blue link 'View complete documentation →'.